

# Predictive Auto-scaling in the Kubernetes Cluster Manager

---

# ■ THANK YOU

- Professor Albrecht, Professor, Williams College
- Brendan Burns, Lead Engineer for Kubernetes, Google

1.

# Goals

---

Why do we care?



# General Goal

---

Contribute to distributed system's ability to reliably and resourcefully do large, varying amounts of computation.

# ■ A **typical** use case

Seek to reliably and resourcefully serve **vote**facts.com until the next election...



# Accomplishing General Goals

---

How do cluster managers reliably and resourcefully perform large, varying amounts of work?

# ■ What is a cluster manager?

- A **cluster** is a collection of commodity computers linked by a local-area network.
- A **cluster manager** admits/runs/monitors user submitted jobs on the cluster.



# Benefits of **Cluster Managers**



Cluster managers allow us to perform computational work that could never be performed on a single computer.



# ■ What are some **cluster managers**?

## **Borg**

Decades old cluster manager from Google. The closed-source precursor to Kubernetes.

## **Mesos**

A low-level cluster manager. If Borg is Ubuntu, Mesos is the Linux Kernel.

## **Apache YARN**

A cluster manager originally for Apache Hadoop.



# Kubernetes





# Specific Goal

---

To maximize the sum of two Kubernetes's metrics: Efficient Resource Utilization and Quality of Service

# Unpacking this Goal

## Kubernetes

An open-source cluster manager from Google.

## Efficient Resource Utilization (ERU)

Is the application efficiently using the resources it is given?

## Quality of Service (QOS)

Is the application accomplishing its stated purpose?

# The goal is **balance.**

---

Increasing ERU/QOS while decreasing the other is easy, we seek to increase the summation.

# Kubernetes specific terms

## Pods

A stateless, replicable wrapper around related containerized applications (ex. a pod for [votefacts.com](https://vote.facts.com) contains an Apache web server and a cache)


## Replication Controllers

A controller for ensuring a given number of replica pods exist.

## Services

A single point of load-balancing access for requests to replica pods.

# Architecture



# Accomplishing Specific Goals

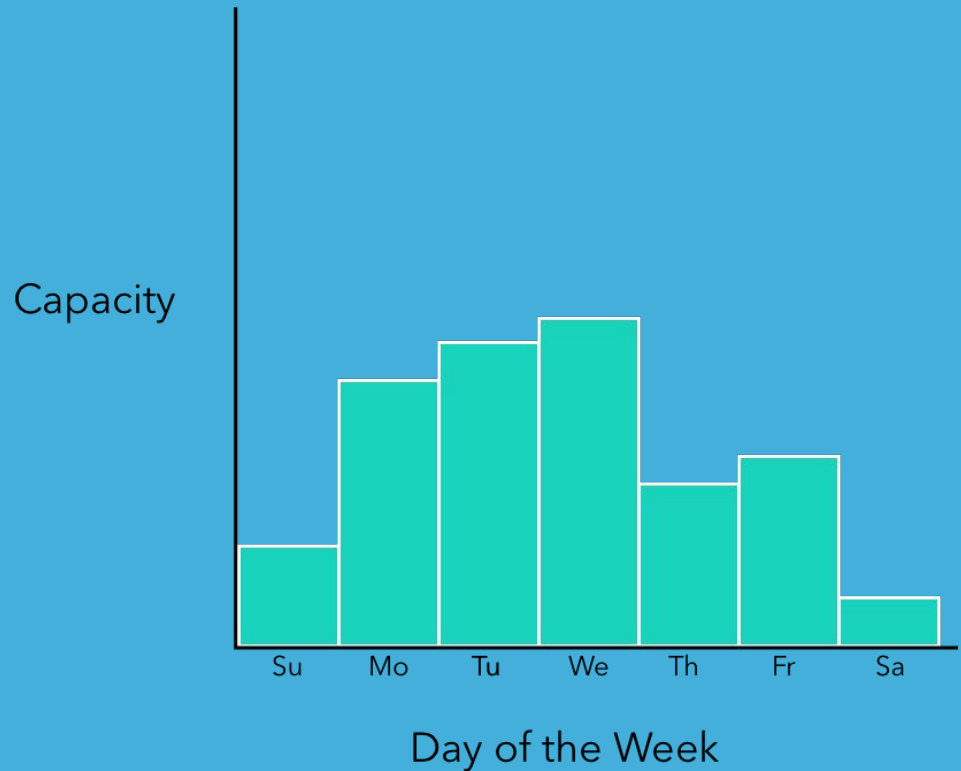
---

How does (predictive) auto-scaling in Kubernetes improve the summation of ERU and QOS?



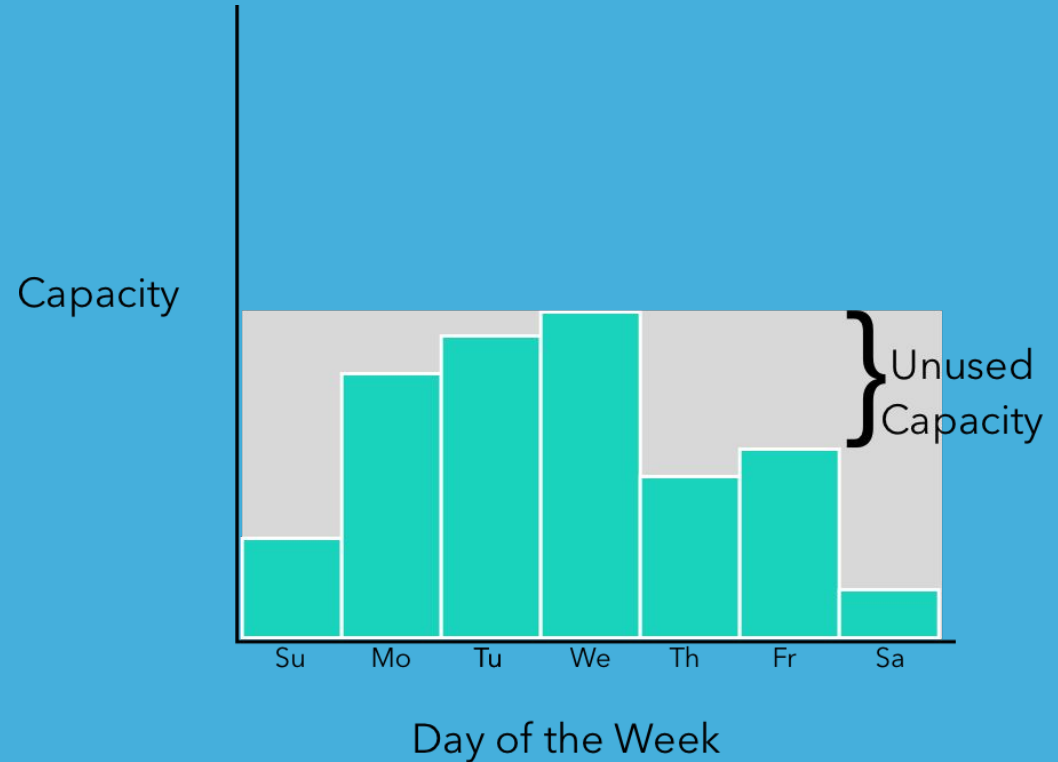
# Benefits of auto-scaling

Imagine the following capacity for [votefacts.com](#) when running on a cluster manager...



# If we do **not** have auto-scaling


We can assign our application the most resources it will ever need... but poor ERU.



# If we do **not** have auto-scaling

We can assign the average amount of capacity needed, improving ERU, but decreasing QOS.

Capacity




# If we do have auto-scaling

We can assign the application the exact resources it needs, when it needs them...

improving the summation of ERU and QOS.

Capacity



# What are the **different characteristics** of auto-scaling?

## Horizontal vs Vertical

How is an application given the extra resources that it needs?

## Reactive vs Predictive

Does auto-scaling occur based on the current or future state of the cluster?

# What are the **major types** of auto-scaling?

## **Threshold-based Rule Policies**

Scale if the current resource usage is not in accordance with a set of predefined rules.

## **Time-series Analysis**

Auto-scale based on repeating pattern in the application load.

## **Control-theory (Feedback Control)**

Scale such that the resource usage is in accordance with predefined guidelines.



# Current State of **Auto-scaling** in Kubernetes

---

Kubernetes currently implements reactive,  
horizontal feedback control based auto-  
scaling of pods.

# Concerns with Auto-scaling in Kubernetes

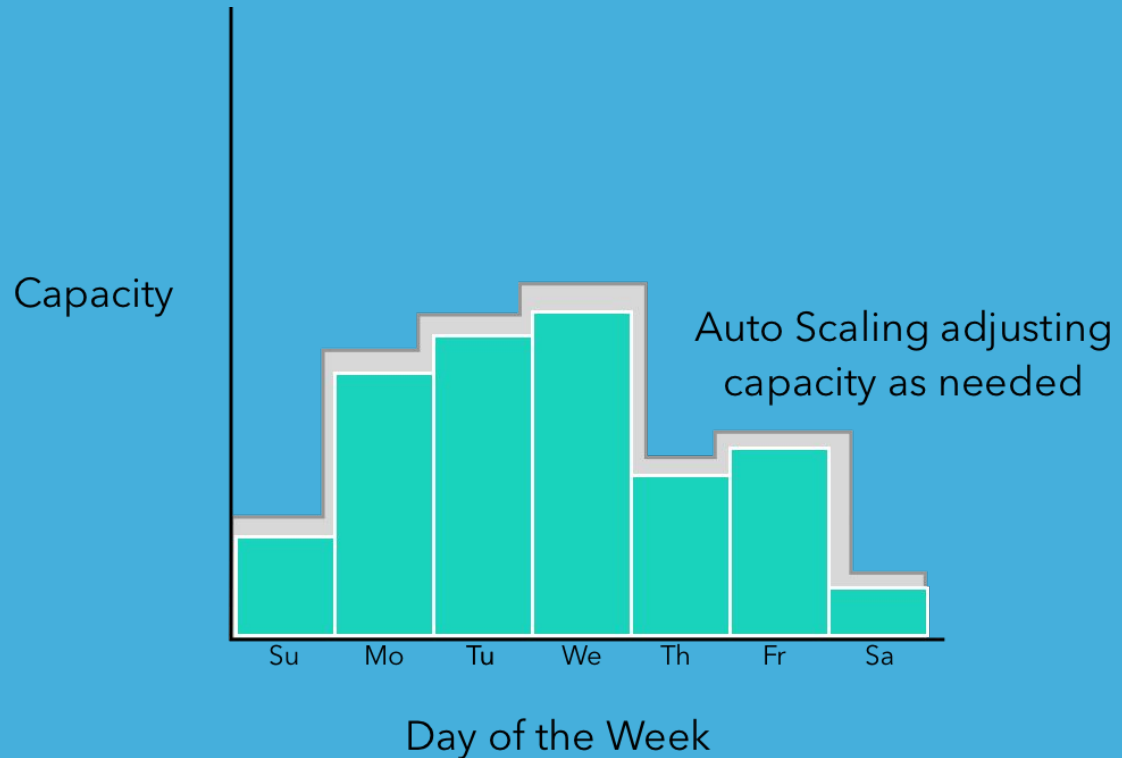
---

Are there ways to improve the summation of ERU and QOS?



# Delayed Pod Initialization Time

What if it takes a long time for a pod to be ready to share in the computational work?



# Improvements to Auto-scaling in Kubernetes

---

What if we add prediction?

## ■ Benefits of adding prediction

- Predictive, horizontal feedback control based auto-scaling of pods
- Improves QOS without decreasing ERU

## ■ A case study

- Imagine at 5:50pm, [vote-facts.com](https://vote-facts.com) needs 100 pods, and at 6pm election results are released, so we need 200 pods.
- Imagine pods take 10 minutes to download all of the election data and initialize.

# Reactive



For **10 minutes**, votefacts.com operates with only **half** the resources it needs, while we wait for the replica pods to initialize.



# Predictive



**votefacts.com** **always** has the resources that it needs.



# ■ Implementation questions?

How long does it take  
for a pod to be ready  
to share in the work?

How can we predict  
the future resource  
utilization of an  
application?

Should this behavior  
be enabled by  
default?


# Status of Work

---


What has been done and what is left to do?



# Current State



# Future Work



# Evaluation

---

How will we know if we're successful?

# Does predictive auto-scaling **increase** ERU + QOS?

How to combine ERU  
and QOS?

What applications  
will we try to auto-  
scale?

What will be the  
external environment  
of these applications?

# THANKS!

Any questions?

---

# CREDITS and CITATIONS

Special thanks to all the people who made and released these awesome resources for free:

- Presentation template by SlidesCarnival
- Photographs by Unsplash
- Thanks to Andrew Udell for assistance with the graphs.
- All Kubernetes info is from <http://kubernetes.io/>.
- Lorigo-Botrán, T., Miguel-Alonso, J., and Lozano, J. A. Auto-scaling Techniques for Elastic Applications in Cloud Environments. Research EHU-KAT-IK, Department of Computer Architecture and Technology, UPV/EHU, 2012.